

## Module 3

**Q. List and explain the core business drivers behind the NoSQL movement.**

=>

### 1. Introduction

- NoSQL (Not Only SQL) refers to a family of non-relational databases designed to overcome the limitations of traditional relational database management systems (RDBMS).
- The movement emerged as businesses began generating massive, diverse, and fast-changing data that traditional systems struggled to handle efficiently.
- Several business and technical factors have driven organizations to adopt NoSQL solutions.

### 2. Core Business Drivers

#### 1. Need for Scalability

- **Problem with RDBMS:**
  - Traditional SQL databases scale vertically (adding more CPU, RAM to a single server).
  - Vertical scaling is expensive and has physical limits.
- **NoSQL Solution:**
  - NoSQL databases scale horizontally (adding more commodity servers to share the load).
  - Distributed systems like Cassandra or MongoDB handle billions of records efficiently.
- **Business Impact:**
  - Enables applications like social media platforms, e-commerce, and IoT systems to serve millions of users simultaneously.

#### 2. Handling Big Data and High Velocity

- **Problem with RDBMS:**

- Structured tables cannot efficiently manage huge, unstructured datasets (videos, images, logs, sensor data).
- High ingestion rates cause performance bottlenecks.

- **NoSQL Solution:**

- Designed to store structured, semi-structured, and unstructured data without rigid schemas.
- Databases like Hadoop HBase and Couchbase can handle high-velocity streaming data.

- **Business Impact:**

- Companies can capture and process real-time data for analytics, fraud detection, and recommendations.

### **3. Flexibility and Schema-less Design**

- **Problem with RDBMS:**

- Schema changes (adding new columns or modifying tables) require downtime and migration, which slows development.

- **NoSQL Solution:**

- Schema-less or dynamic schemas allow storing different attributes in different records.
- Document-oriented databases like MongoDB let developers add new fields on the fly without breaking existing data.

- **Business Impact:**

- Faster application development and iteration, critical for startups and agile teams.

### **4. Cloud and Distributed Computing Adoption**

- **Problem with RDBMS:**

- Traditional databases are not inherently cloud-friendly and require complex manual partitioning.

- **NoSQL Solution:**

- Designed for distributed architectures across multiple nodes and data centers.
- Built-in replication and partitioning (sharding) make them naturally suitable for cloud platforms.

- **Business Impact:**

- Organizations reduce infrastructure costs and gain high availability and geographical distribution.

## **5. Cost Efficiency Using Commodity Hardware**

- **Problem with RDBMS:**

- Requires high-end servers to support vertical scaling.

- **NoSQL Solution:**

- Runs on low-cost commodity hardware in clusters.
- Automatic failover and replication avoid the need for expensive hardware redundancy.

- **Business Impact:**

- Significant cost savings while still achieving high performance and reliability.

## **6. Demand for High Availability and Fault Tolerance**

- **Problem with RDBMS:**

- Centralized architecture means single point of failure.

- Complex clustering and replication mechanisms are required for availability.
- **NoSQL Solution:**
  - Provides automatic replication and eventual consistency across nodes.
  - Systems like Cassandra guarantee always-on service with no downtime.
- **Business Impact:**
  - Businesses can provide 24/7 services even during node failures or maintenance.

## 7. Support for Modern Application Requirements

- **Problem with RDBMS:**
  - Difficult to handle location-based queries, user personalization, graph relationships, and real-time analytics.
- **NoSQL Solution:**
  - Offers specialized databases:
  - Key-Value stores for caching and session management (Redis).
  - Document stores for content management (MongoDB).
  - Graph databases for social networks (Neo4j).
- **Business Impact:**
  - Enables innovative features like personalized recommendations, fraud detection, and social graph analysis.

## Q. What is a Key-Value Store?

=>

### Definition

- A **key-value store** is the simplest type of NoSQL database where **data is stored as pairs of keys and values**.
- Each **key** is a unique identifier, and its **value** can be any kind of data (string, number, JSON, image, etc.).
- The database retrieves data **by key only**, ensuring **fast lookups** without complex queries.

### How it Works

- Think of it as a **dictionary (hash map)**:
  - Key → "user123"
  - Value → {"name":"Alice", "age":25}
- To fetch the user's data, simply ask for the key "user123" — no need to scan the whole database.

### Examples of Key-Value Stores

- **Redis** – in-memory key-value store used for caching and session management.
- **Amazon DynamoDB** – cloud-based key-value store.
- **Riak / BerkeleyDB / Voldemort** – distributed key-value stores for scalable applications.

### Benefits of Using a Key-Value Store

#### 1. High Performance

- **O(1) lookup time:** Data retrieval is extremely fast because keys are indexed internally.
- Ideal for **real-time applications** like gaming leaderboards, chat apps, and financial systems.

## 2. Simplicity

- Data model is straightforward — just keys and values.
- Easy to understand, implement, and integrate into applications without complex schemas.

## 3. Scalability

- Designed for **horizontal scaling** by adding more nodes.
- Distributed key-value stores (e.g., DynamoDB, Riak) automatically partition and replicate data.
- Handles **massive traffic and big datasets** efficiently.

## 4. Flexibility of Data Types

- Values can be **any type**: text, JSON, binary files, or serialized objects.
- No fixed schema, so developers can **store varied data without altering the structure**.

## 5. Fault Tolerance and High Availability

- Many key-value stores have **built-in replication** to keep copies of data on multiple nodes.
- Even if one node fails, the data is still available, making them highly reliable for critical systems.

## 6. Cost Efficiency

- Works well on **commodity hardware** and cloud-based clusters.
- Minimal operational overhead compared to traditional relational databases.

**Q. What is graph store? Give an example where a graph store can be used to effectively solve a particular business problem**

=>

### Definition

- A **graph database (GDB)** is a type of NoSQL database that stores data using **graph structures** instead of traditional tables or documents.
- Data is represented as:
  - **Nodes:** Entities (e.g., users, products, locations)
  - **Edges:** Relationships between nodes (e.g., friendship, transactions, connections)
  - **Properties:** Attributes of nodes or edges (e.g., name, age, timestamp)
- Graph databases are **designed to efficiently handle relationships** and enable fast queries across connected data.
- **Examples:** Neo4j, Amazon Neptune, ArangoDB

### Graph Representation

- Nodes store **data entities**, and edges capture **relationships** between them.
- Example structure for a social network:

#### Nodes (Users):

id	first name	last name	email	phone
1	Anay	Agarwal	anay@example.net	555-111-5555
2	Bhagya	Kumar	bhagya@example.net	555-222-5555

3	Chaitanya	Nayak	chaitanya@example.net	555-333-5555
---	-----------	-------	-----------------------	--------------

### Edges (Friendship relationships):

user_id	friend_id
3	1
3	2
3	4
3	5

Querying friends of a user in a **graph database** is faster because relationships are **directly stored as edges**.

### When Do We Need a Graph Database?

#### 1. Many-to-Many Relationships

- Ideal for networks where entities are connected to multiple other entities (e.g., social media friends).

#### 2. Relationship-Centric Queries

- When **relationships are more important than individual data**, like tracking user interactions or supply chains.

#### 3. Low Latency on Large Datasets

- Graph databases can retrieve connected data in **constant time ( $O(1)$  per connection)**, unlike relational joins that can become very slow.



## Business Example: Social Network Friend Query

- **Problem:** Chaitanya wants to see all her friends' profiles.

### Relational DB Approach:

- Query involves joining **Users** table and **Friendship** table.
- Time complexity:  $O(M * \log(N))$  for M queries on N friendships.

### Graph DB Approach:

- Locate Chaitanya node, traverse edges to friends.
- Time complexity:  $O(N)$  (single-step traversal per connection).

### Result:

- Faster retrieval, especially as the network grows large.
- Real-time queries become feasible for millions of users.

## Advantages of Graph Databases

- Handles **frequent schema changes** easily.
- Efficient for **managing large volumes of interconnected data**.
- Enables **real-time query responses**.
- Supports **intelligent data activation** like recommendations or fraud detection.

## Disadvantages / Limitations

- Not always the best solution for all applications.
- Horizontal scaling can be challenging; may affect performance.
- Updating all nodes with a given parameter can be inefficient.
- May not outperform other NoSQL options in certain scenarios.

## Q. CAP vs ACID

=>

Feature/Aspect	ACID (Relational DB)	CAP (NoSQL / Distributed DB)
<b>Definition</b>	Guarantees Atomicity, Consistency, Isolation, Durability in transactions	Guarantees Consistency, Availability, Partition tolerance in distributed systems
<b>Primary Focus</b>	Transaction correctness in a single database	System behavior under network partitions in distributed databases
<b>Consistency</b>	Strong consistency: database always remains valid after transactions	Trade-off: can be strong or eventual consistency depending on choice
<b>Availability</b>	High availability depends on the DB setup, but not a core principle	Must choose between Availability or Consistency in presence of network failures
<b>Partition Tolerance</b>	Not explicitly considered; assumes single node or tightly coupled system	Core requirement: system continues to operate even if network partitions occur
<b>Use Cases</b>	Banking, financial systems, ERP — critical transactions	Social media, large-scale web apps, e-commerce — scalable and distributed data
<b>Implementation</b>	Traditional RDBMS (MySQL, PostgreSQL, Oracle)	Distributed NoSQL (Cassandra, MongoDB, DynamoDB)

<b>Transaction Support</b>	Full ACID transactions supported	Often limited or eventual transaction guarantees
<b>Scalability</b>	Vertical scaling (bigger servers)	Horizontal scaling (add more nodes)
<b>Trade-offs</b>	Prioritizes correctness over availability in failures	Must balance C, A, P — cannot achieve all three simultaneously

**Q. List the architectural patterns in NoSQL databases. Discuss the Key-Value and Document-Oriented patterns, focusing on their characteristics, use cases, and examples.**

=>

## **1. Introduction**

- NoSQL databases are designed to handle **large-scale, distributed, and unstructured data**.
- Unlike traditional relational databases, NoSQL databases follow **different architectural patterns** depending on the type of data and use case.
- Common architectural patterns include:
  1. **Key-Value Stores**
  2. **Document-Oriented Stores**
  3. **Column-Family Stores**
  4. **Graph Databases**

## **2. Key-Value Store Pattern**

### **Definition**

- Data is stored as **key-value pairs**: each key is unique, and the value can be **any type of data** (text, JSON, binary, etc.).
- The system retrieves data **by key only**, providing **fast lookups**.

### **Characteristics**

- Extremely **simple data model**.
- Highly **scalable** via horizontal partitioning (sharding).
- Designed for **high performance and low-latency reads/writes**.
- Schema-less: no fixed structure for values.

## Use Cases

- Caching (e.g., Redis for session data).
- Real-time analytics.
- User profiles, preferences, or settings storage.
- IoT or sensor data ingestion.

## Examples

- Redis, DynamoDB, Riak, Voldemort.

## 3. Document-Oriented Store Pattern

### Definition

- Data is stored as **documents**, often in **JSON, BSON, or XML format**.
- Each document contains **self-describing data**, often including nested structures.
- Documents are grouped into **collections**, and each document has a unique key (document ID).

### Characteristics

- Supports **flexible, dynamic schemas**.
- Can query based on **document fields** rather than just key.
- Handles **complex data structures** naturally.
- Horizontally scalable through **sharding** and **replication**.

### Use Cases

- Content management systems (CMS) storing articles or blog posts.

- E-commerce platforms storing product catalogs.
- Event logging and analytics with semi-structured data.
- Applications requiring **rapid iteration and schema evolution**.

## Examples

- MongoDB, CouchDB, ArangoDB, Amazon DocumentDB.

## 4. Comparison of Key-Value vs Document-Oriented Patterns

Feature	Key-Value Store	Document-Oriented Store
<b>Data Model</b>	Simple key → value mapping	Documents with fields and nested structures
<b>Query Flexibility</b>	Only by key	By key and document fields
<b>Schema</b>	Schema-less, unstructured	Flexible schema, supports complex/nested data
<b>Use Case</b>	Caching, real-time session storage	CMS, product catalogs, event logs
<b>Performance</b>	Extremely fast, low latency	Fast, slightly more complex queries
<b>Examples</b>	Redis, DynamoDB, Riak	MongoDB, CouchDB, ArangoDB

**Q. Describe the four ways by which big data problems are handled by NoSQL.**

**=>**

- Big data is characterized by **Volume, Velocity, Variety, and Veracity**.
- Traditional relational databases often struggle with these challenges due to **rigid schemas, limited scalability, and high latency**.
- **NoSQL databases** address big data problems through **innovative architectural approaches** that enable distributed storage, flexible schemas, and high performance.

## **2. Four Ways Big Data Problems Are Handled by NoSQL**

### **1. Horizontal Scalability**

- **Definition:** Distributing data across multiple nodes (servers) instead of relying on a single powerful machine.
- **How it helps:**
  - Handles **huge volumes of data** by adding more nodes to the cluster.
  - Reduces performance bottlenecks by parallelizing reads and writes.
- **Example:**
  - **Cassandra** automatically partitions data across multiple nodes for linear scalability.
- **Benefit:** Avoids the cost and limits of vertical scaling in relational databases.

### **2. Schema Flexibility**

- **Definition:** No fixed schema; data structure can evolve without downtime.
- **How it helps:**
  - Handles **variety of data**: structured, semi-structured (JSON/XML), or unstructured (images, logs).

- Allows **rapid development and iterative changes** without complex migrations.
- **Example:**
  - **MongoDB** stores JSON-like documents with varying fields.
- **Benefit:** Supports fast-changing business requirements and diverse datasets.

### 3. High Availability and Fault Tolerance

- **Definition:** Ensures data is accessible even when nodes fail.
- **How it helps:**
  - Uses **replication** to keep multiple copies of data across nodes.
  - Provides **continuous availability** for real-time applications.
- **Example:**
  - **Riak** replicates data across nodes so queries succeed even if some nodes are down.
- **Benefit:** Enables robust systems for critical applications like e-commerce, social media, or banking.

### 4. Optimized for High-Velocity Data

- **Definition:** Efficient ingestion and retrieval of rapidly changing or streaming data.
- **How it helps:**
  - Handles **millions of writes per second** without slowing down the system.
  - Supports **real-time analytics** and operational monitoring.
- **Example:**
  - **Redis** and **Cassandra** can process real-time logs or sensor data at scale.
- **Benefit:** Supports modern applications like IoT, clickstream analysis, and recommendation engines.